

Buffer overflows

JOÃO PAULO BARRACA

BO - According to CAPEC-100

- **Targets improper or missing bounds checking on buffer operations**
 - typically triggered by input injected by an adversary.
- **An adversary is able to write past the boundaries of allocated buffer regions in memory**
- **Causes a program crash or potentially redirection of execution as per the adversaries' choice.**
 - Denial of Service
 - (Remote) Code Execution

BO - Scope

- **CWE-119 is extremely broad as there are many types of BO**
- **Characteristics of a BO**
 - Type of access: Read or Write
 - Type of memory: stack, heap
 - Location: before or after the buffer
 - Reason: iteration, copy, pointer arithmetic, memory clear, mapping

Other Direct Child CWEs

CWE-120	Buffer Copy without Checking Size of Input ('Classic Buffer Overflow')
CWE-125	Out-of-bounds Read
CWE-466	Return of Pointer Value Outside of Expected Range
CWE-786	Access of Memory Location Before Start of Buffer
CWE-787	Out-of-bounds Write
CWE-788	Access of Memory Location After End of Buffer
CWE-805	Buffer Access with Incorrect Length Value
CWE-822	Untrusted Pointer Dereference
CWE-823	Use of Out-of-range Pointer Offset
CWE-824	Access of Uninitialized Pointer
CWE-825	Expired Pointer Dereference

Relevant CWEs with specific types

CWE-120: Classic Buffer Overflow: copy without checking the size of the input

CWE-121: Stack-based Buffer Overflow: overwrite over data in the Stack Segment

CWE-122: Heap-based Buffer Overflow: overwrite over data in the Heap Segment


CWE-123: Write-what-where Condition: ability to write to any memory of choice

CWE-124: Buffer Underwrite ('Buffer Underflow'): Write to memory before the buffer

CWE-126: Buffer Over-read: Read after the buffer ends (e.g., using an index)

CWE-127: Buffer Under-read: Read before the buffer start (e.g., using an index)

Total Matches By Year



Popularity at NVD

Popularity decline

- **Better tools to check for the vulnerability**
 - Static/Dynamic Code analysis
- **Dissemination of bound checking mechanisms in compilers**
 - Standard in most distributions and enabled by default
 - Still lacking in embedded devices
- **Increasingly higher adoption of higher layer languages**
 - Extensive use and Open Sources libraries improves security
 - Security focused languages such as Rust

Potentially Vulnerable Software

- **Any software that gets information from external sources**
 - Sockets, PIPEs and other IPC
 - Files
 - Program arguments
 - Environment Variables
- **Software developed in languages with direct memory access**
 - Mostly C and C++ (or at least with most devastating impact)
 - But also: Go when using “unsafe”, PHP, Python, Java, etc...

Dominant prevalence

- **Anything that was made in a language with access to memory**
 - Server software packages (nginx, apache, mysql, ...)
- **Embedded and IoT devices**
 - Due to lack of compiler support
 - Due to lack of hardware capabilities

... in python

```
# bo_1.py

message = "Hello World"

buffer = [None] * 10

print(message)

for i in range(15):

    buffer[i] = 'A'

print(message)
```

```
$ python3 bo_1.py

Hello World

Traceback (most recent call last):

  File "bo_1.py", line 7, in <module>

    buffer[i] = 'A'

IndexError: list assignment index out of
range
```

... in C

```
#include <stdio.h>

void main(int argc, char* argv){
    char message[] = "Hello World";
    int buffer[5];
    int i;

    printf("%s\n", message);
    for(i = 0; i < 15; i++) {
        buffer[i] = 'A';
    }
    printf("%s\n", message);
}
```

```
./bo_1

Hello World

AAAAAAAAAAAAAAAAAd AAAAAAAAAAd
```

Vulnerabilities in languages (mostly C/C++)

- **Not memory safe: programmers can read/write memory freely and are not constrained by the address or size of the variables**
 - Great flexibility, but huge risk as mistakes lead to accessing memory that otherwise should not be accessed
 - C/C++ compilers have freedom to optimize code and even sometimes undefined behavior
- **Memory safe languages intercept such errors, raising errors**
 - Program will crash (DoS), but impact is limited

```
// Correct usage
printf("%d\n", *value);

// Reading memory after the variable
printf("%d\n", *(value + 4));

// Reading memory before the variable
printf("%d\n", *(value - 4));
```

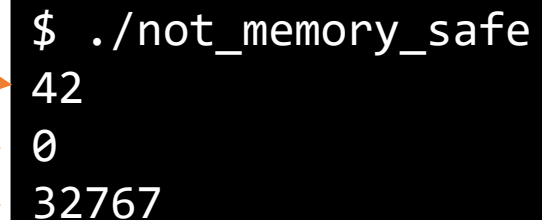
Vulnerabilities in languages (mostly C/C++)

- **Not memory safe: programmers can read/write memory freely and are not constrained by the address or size of the variables**
 - Great flexibility, but huge risk as mistakes lead to accessing memory that otherwise should not be accessed
 - C/C++ compilers have freedom to optimize code and even sometimes undefined behavior
- **Memory safe languages intercept such errors, raising errors**
 - Program will crash (DoS), but impact is limited

```
// Correct usage
printf("%d\n", *value);

// Reading memory after the variable
printf("%d\n", *(value + 4));

// Reading memory before the variable
printf("%d\n", *(value - 4));
```



```
$ ./not_memory_safe
42
0
32767
```

Vulnerabilities in languages (mostly C/C++)

- **Not type safe: memory content can be reinterpreted as required by the programmer**
 - Casts may be arbitrarily allowed and not checked
- **Type safe languages do not allow reinterpretation, or only safe reinterpretation**
 - Cast a byte to int is safe, a buffer to int is not.

```
int value = 42;

// Correct usage
printf("%d\n", value);

// Cast to variable with different storage
printf("%f\n", *((double*) &value));

// Cast to variable with different size
printf("%llu\n", *((unsigned long long*) &value));
```

Vulnerabilities in languages (mostly C/C++)

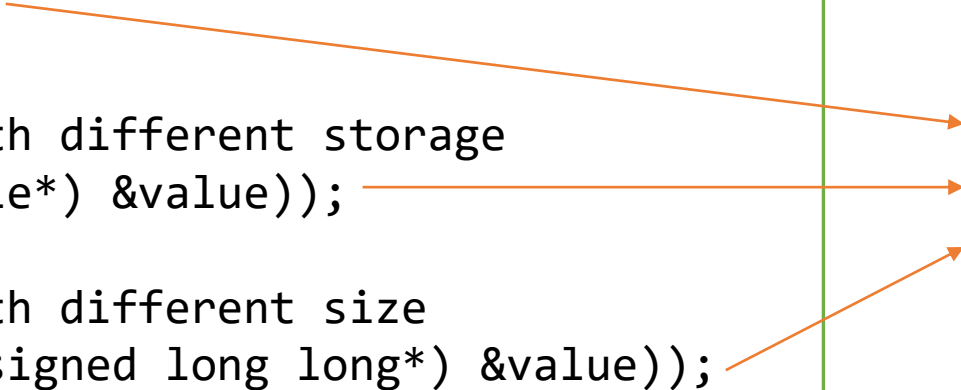
- **Not type safe: memory content can be reinterpreted as required by the programmer**
 - Casts may be arbitrarily allowed and not checked
- **Type safe languages do not allow reinterpretation, or only safe reinterpretation**
 - Cast a byte to int is safe, a buffer to int is not.

```
int value = 42;

// Correct usage
printf("%d\n", value);

// Cast to variable with different storage
printf("%f\n", *((double*) &value));

// Cast to variable with different size
printf("%llu\n", *((unsigned long long*) &value));
```



```
$ ./not_type_safe
42
0.000000
1170988679674462250
```

Vulnerabilities in languages (mostly C/C++)

➤ Dynamically allocated memory has no implicit management mechanism

- Programmer must allocate and deallocate all memory
- Programmer must know how memory was allocated
- Programmer must free memory only after there is no other reference

```
char* buffer = (char*) malloc(10);  
char* str = buffer;  
  
free(buffer);  
  
// Write after free (and write beyond buffer)  
memcpy(str, "Hello World!!!!", 15);  
// Read after free (and read beyond buffer)  
printf("%s\n", str);
```

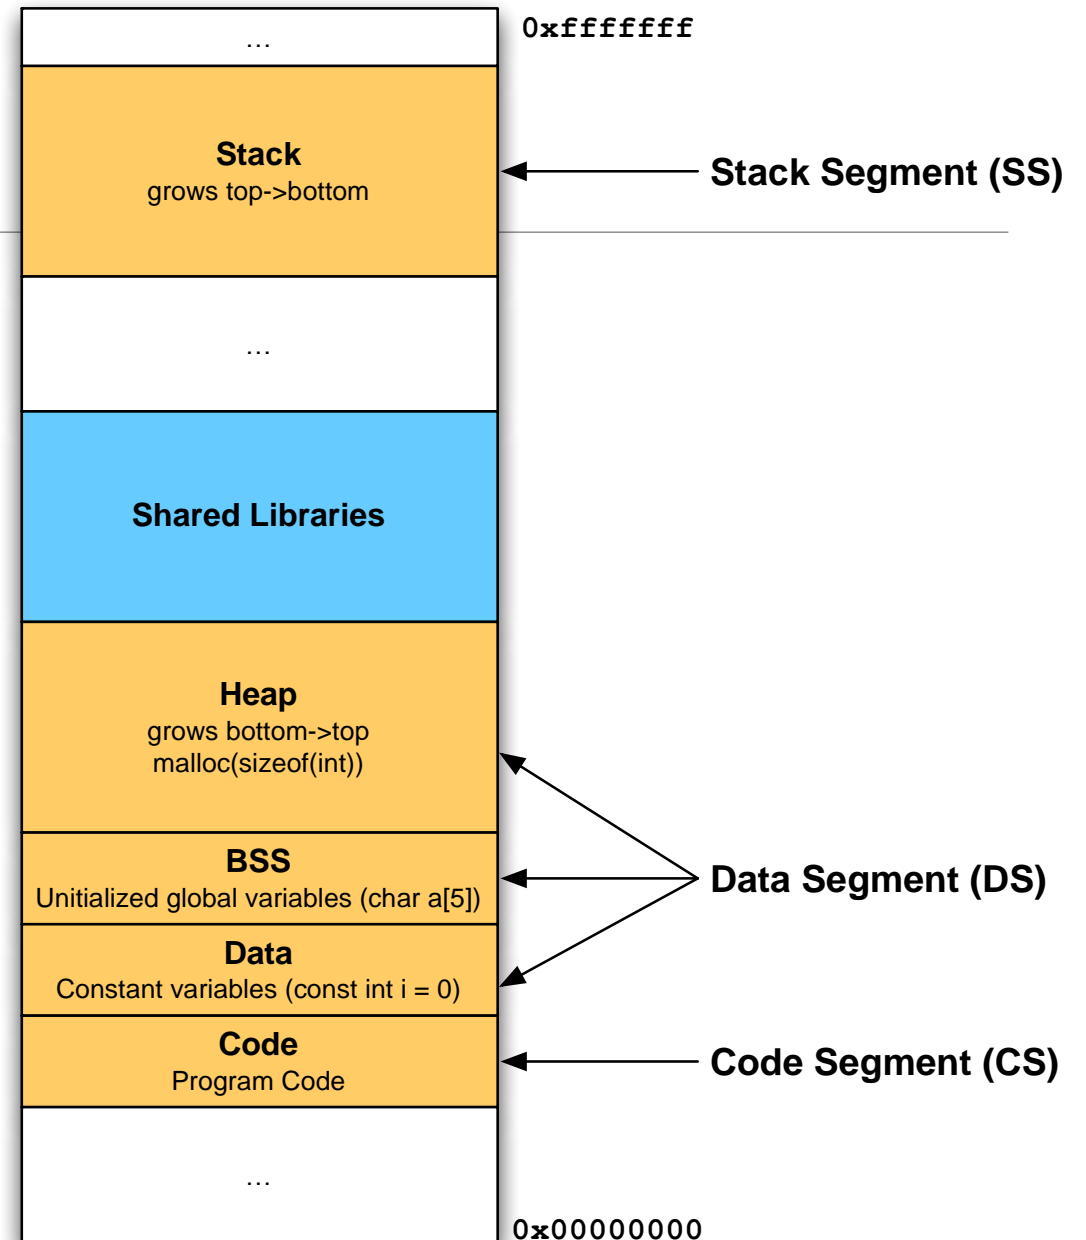
```
$ ./dynamic_memory  
Hello World!!!!
```


Why? Memory Structure 101

- **Kernel organizes memory in pages**
 - Typically 4096 bytes
- **Processes operate in a Virtual Memory Space**
 - Mapped to real pages, which can be in RAM or Swapped
- **Kernel splits program in several segments**
 - Increases security
 - segment based permissions
 - Increases performance
 - some are dynamic: invalidated when program terminates
 - some are static: can be retained, speed repeated startup

Memory Structure

- **SS: Local variables and execution flow**
- **Shared Libraries: .so/dlls loaded.**
 - Addresses are shared between programs
- **Heap: memory allocated with malloc/new**
- **BSS: Global Variables**
- **Data: Constants**
- **Code: Actual instructions**



mem.c (available in course web page)

➤ Simple program showing the memory map of itself

➤ Features:

- Prints the address of objects of different types
 - Argument
 - Dynamic memory with malloc
 - Global Variable
 - Constant
 - Function
- Prints the memory maps as exposed in /proc/self/maps
- Creates a recursive function and prints the address of local variables
- Crashes with a Stack Overflow

mem.c

Internal Variables (Page = 4096)


`&argc = bfeb8590 -> stack = bfeb8000`

`malloc = 08435008 -> heap = 08435000`

`bssvar = 0804a034 -> bss = 0804a000`

`cntvar = 08048920 -> const = 08048000`


`&main = 0804865c -> text = 08048000`



mem.c

Content of /proc/self/maps

```
08048000-08049000 r-xp 00000000 08:01 26845750 /home/s/mem
08049000-0804a000 r--p 00000000 08:01 26845750 /home/s/mem
0804a000-0804b000 rw-p 00001000 08:01 26845750 /home/s/mem
08435000-08456000 rw-p 00000000 00:00 0 [heap]
b7616000-b7617000 rw-p 00000000 00:00 0
b7617000-b776a000 r-xp 00000000 08:01 1574823 /lib/tls/i686/cmov/libc-2.11.1.so
b776a000-b776b000 ---p 00153000 08:01 1574823 /lib/tls/i686/cmov/libc-2.11.1.so
b776b000-b776d000 r--p 00153000 08:01 1574823 /lib/tls/i686/cmov/libc-2.11.1.so
b776d000-b776e000 rw-p 00155000 08:01 1574823 /lib/tls/i686/cmov/libc-2.11.1.so
b776e000-b7771000 rw-p 00000000 00:00 0
b777e000-b7782000 rw-p 00000000 00:00 0
b7782000-b7783000 r-xp 00000000 00:00 0 [vdso]
b7783000-b779e000 r-xp 00000000 08:01 1565567 /lib/ld-2.11.1.so
b779e000-b779f000 r--p 0001a000 08:01 1565567 /lib/ld-2.11.1.so
b779f000-b77a0000 rw-p 0001b000 08:01 1565567 /lib/ld-2.11.1.so
bfe99000-bfeba000 rw-p 00000000 00:00 0 [stack]
```




mem.c

Stack evolution:

```
foo [000]: &argc = bfeb8140 -> stack = bfeb8000
foo [001]: &argc = bfdb8110 -> stack = bfdb8000
foo [002]: &argc = bfc8b80e0 -> stack = bfc8b8000
foo [003]: &argc = bfbb80b0 -> stack = bfbb8000
foo [004]: &argc = bfab8080 -> stack = bfab8000
foo [005]: &argc = bf9b8050 -> stack = bf9b8000
foo [006]: &argc = bf8b8020 -> stack = bf8b8000
foo [007]: &argc = bf7b7ff0 -> stack = bf7b7000
foo [008]: &argc = bf6b7fc0 -> stack = bf6b7000
```

Segmentation fault



Stack organization

- **Stack is organized by frames, one for each function call**
 - Memory reserved for the function to use as it requires
- **Each stack frame stores:**
 - Return Information
 - Local Variables
 - Arguments to following functions (x32: all, x64: +5th)

```
void main(){  
    foo();  
}
```

```
void foo(){  
    bar();  
}
```



Stack organization

- **Stack is organized by frames, one for each function call**
 - Memory reserved for the function to use as it requires
- **Each stack frame stores:**
 - Return Information
 - Local Variables
 - Arguments to following functions



Stack organization

- **Return information has 2 major objectives**
 - Chaining frames as new functions are called
 - Return to the next instruction after the function ends
- **Frame chaining**
 - When a function is called, the address of the current stack frame (Register RBP in x64) is push to the frame
 - When the function ends, RBP is popped
 - Caller function has it's frame restored
- **Function chaining**
 - When a function is called, the address of the next instruction is push to the stack (RIP register)
 - When a function ends, that address is popped
 - Execution resumes at the caller function



mem_local.c (available in course web page)

➤ Prints the address to several variables

- Local variables declared in the main function
- Arguments passed to the foo function
- Local variables in the foo function

main

argc : 0x7fffd6baeddc
argv : 0x7fffd6baeed8

foo

a : 0x7fffd6baed8c
local_a: 0x7fffd6baed9b
buffer : 0x7fffd6baeda0
local_b: 0x7fffd6baed9c

```
char foo(int a,){
    char local_a = 3;
    char buffer[16];
    int local_b = 5;

    printf("%p\n", &a);
    printf("%p\n", &local_a);
    printf("%p\n", &buffer);
    printf("%p\n", &local_b);

    buffer[0] = local_a;
    return buffer[0];
}

int main(int argc, char* argv[]){
    printf("%p\n", &argc);
    printf("%p\n", argv);

    return foo(argc);
}
```

mem_local.c – Conclusions

➤ Stack frame grows from higher addresses to lower addresses

- Main has variables at 0xbaedb.
- Foo has variables at 0xbaed6-8.

main

argc : 0x7fffd6baeddc
argv : 0x7fffd6baeed8

foo

a : 0x7fffd6baed8c
local_a: 0x7fffd6baed9b
buffer : 0x7fffd6baeda0
local_b: 0x7fffd6baed9c

```
char foo(int a,){
    char local_a = 3;
    char buffer[16];
    int local_b = 5;

    printf("%p\n", &a);
    printf("%p\n", &local_a);
    printf("%p\n", &buffer);
    printf("%p\n", &local_b);

    buffer[0] = local_a;
    return buffer[0];
}

int main(int argc, char* argv[]){
    printf("%p\n", &argc);
    printf("%p\n", argv);

    return foo(argc);
}
```

mem_local.c – Conclusions

- Declaration order doesn't matter!
- Compiler will place variables where he seems adequate
 - Will keep information aligned
 - May create empty spaces
 - May deploy additional protection mechanisms (canaries)

main

argc : 0x7fffd6baeddc
argv : 0x7fffd6baeed8

foo

a : 0x7fffd6baed8c
local_a: 0x7fffd6baed9b (3rd)
buffer : 0x7fffd6baeda0 (1st)
local_b: 0x7fffd6baed9c (2nd)

```
char foo(int a,){  
    char local_a = 3;  
    char buffer[16];  
    int local_b = 5;  
  
    printf("%p\n", &a);  
    printf("%p\n", &local_a);  
    printf("%p\n", &buffer);  
    printf("%p\n", &local_b);
```

```
    buffer[0] = local_a;  
    return buffer[0];
```

```
}
```


```
int main(int argc, char* argv[]){  
    printf("%p\n", &argc);  
    printf("%p\n", argv);
```

```
    return foo(argc);
```

```
}
```

mem.c

Q: How much can it grow?




mem.c

1. Until a limit imposed by the SO is reached. Ex:

- glibc i386, x86_64 7.4 MB
- Tru64 5.1 5.2 MB
- Cygwin 1.8 MB
- Solaris 7..10 1 MB
- MacOS X 10.5 460 KB
- AIX 5 98 KB
- OpenBSD 4.0 64 KB
- HP-UX 11 16 KB

2. Until vital memory is overwritten

- ...mostly in embedded devices



CWE-120 Classic Overflow

- **Given an input buffer, data is copied without checking its size**
 - If destination buffer is larger than input data, nothing bad happens
 - If destination buffer is smaller than input data, memory is overwritten

- **Impact: memory is overwritten**
 - Mostly affects local variables
 - May change the execution flow
 - Change of local control variables
 - Change of stored Instruction Pointer
 - May be used to inject external code

- **Solution: take in consideration the size of the destination buffer!**

Classic Overflow – prog 1

➤ Description:

- Reads the username from the command line
- Input is stored in variable **username**
- Variable can hold strings up to 31 chars
 - Why 31 and not 32?
- **gets** functions has no limit on input size
- **printf** will print the content

➤ Shows a simple write beyond boundaries

- printf also shows a read beyond boundaries

```
//classic/prog_1.c
//gcc -O0 -fno-stack-protector -o prog_1 prog_1.c

#include <stdio.h>

int main() {
    char username[32];
    puts("username:");
    gets(username);
    printf("Welcome %s!\n", username);
    return 0;
}
```


Classic Overflow – prog 1

➤ Reading more than 31 chars will result in overwriting the memory after the username

- There are no other variables, so this will be stack structures (addressed later)

➤ printf will print chars up to 0x00, potentially printing program memory

- Function is insecure as there are no explicit boundaries except the actual string content

```
//classic/prog_1.c
//gcc -O0 -fno-stack-protector -o prog_1 prog_1.c

#include <stdio.h>

int main() {
    char username[32];
    puts("username:");
    gets(username);
    printf("Welcome %s!\n", username);
    return 0;
}
```

Exercise: classic/prog 1

- Install gef: `pip3 install --user gdb-gef`
- Compile the binary: `gcc -g -O0 -fno-stack-protector -o prog_1 prog_1.c`
- Analyze the execution with different payloads
 - Print register: `p $rsp` or variable address `p &username`
 - Check stack information: `info frame`
- Determine
 - What is the stack base address?
 - Where is the return information?
 - How many bytes can be entered without overflow?
 - How many bytes can be written without damage?
 - What happens when an overflow is achieved?

```

0x00007ffffffdf20 +0x0000: 0x0000000000000000 ← $rsp
0x00007ffffffdf28 +0x0008: 0x000000000000401090 → <_start+0> endbr64
0x00007ffffffdf30 +0x0010: 0x00007fffffffee030 → 0x0000000000000001
0x00007ffffffdf38 +0x0018: 0x0000000000000000
0x00007ffffffdf40 +0x0020: 0x0000000000000000 ← $rbp
0x00007ffffffdf48 +0x0028: 0x00007ffffff5c70b3 → <__libc_start_main+243> mov edi, eax
0x00007ffffffdf50 +0x0030: 0x00007ffffff7dd620 → 0x00030d0b00000000
0x00007ffffffdf58 +0x0038: 0x00007fffffffee038 → 0x00007fffffffee28f

```

```

0x40117a <main+4>      push    rbp
0x40117b <main+5>      mov     rbp, rsp
0x40117e <main+8>      sub     rsp, 0x20
→ 0x401182 <main+12>   lea     rdi, [rip+0xe7b]      # 0x402004
0x401189 <main+19>   call   0x401060 <puts@plt>
0x40118e <main+24>   lea     rax, [rbp-0x20]
0x401192 <main+28>   mov     rdi, rax
0x401195 <main+31>   mov     eax, 0x0
0x40119a <main+36>   call   0x401080 <gets@plt>

```

```

1  #include <stdio.h>
2
3  int main() {
4      char username[32];
5      // username=0x00007ffffffdf20 → 0x0000000000000000
→ 5      puts("username:");
6      gets(username);
7      printf("Welcome %s!\n", username);
8      return 0;
9  }
10

```

[#0] Id 1, Name: "prog_1", stopped 0x401182 in main (), reason: SINGLE STEP

[#0] 0x401182 → main()

Saved \$BP
Saved \$PC

```
0x00007ffffffdf20 +0x0000: "aaaaaaaaaaaaaaaaaaaaaaaaaaaa" ← $rax, $rsp, $r8
0x00007ffffffdf28 +0x0008: "aaaaaaaaaaaaaaaaaaaaaaaaaaaa"
0x00007ffffffdf30 +0x0010: "aaaaaaaaaaaaaaaaaaaa"
0x00007ffffffdf38 +0x0018: 0x0061616161616161 ("aaaaaaa"? )
0x00007ffffffdf40 +0x0020: 0x0000000000000000 ← $rbp
0x00007ffffffdf48 +0x0028: 0x00007ffffff5c70b3 → <__libc_start_main+243> mov edi, eax
0x00007ffffffdf50 +0x0030: 0x00007ffffff7dd620 → 0x00030d0b00000000
0x00007ffffffdf58 +0x0038: 0x00007ffffffee038 → 0x00007ffffffee28f
```

➤ What is the stack base address?

- info frame: 0x7ffffffdf50
- p \$rbp: 0x7ffffffdf40

➤ Where is the return information?

- Just before \$rbp

➤ How many bytes can be entered without overflow?

- sizeof(username) - 1

➤ How many bytes can be written without damage?

- 32
- It could have been different due to empty space

➤ What happens when an overflow is achieved?

- Saved \$BP is overwritten and then Saved \$PC is overwritten
- In this case, 31 'a' were provided and an additional \0 was added at .. edf38

Classic Overflow – classic/prog_2

➤ Flow:

- Asks for username and password
- Validates credentials
- Asks for message
- If user authenticated, access is granted

➤ Issues:

- Several uncontrolled reads
- All variables may overwrite other

➤ Demonstrates overwrite of local variables

- Each vulnerable variable may overwrite others above

```
int main() {
    char allowed = 0;
    char password[8];
    char username[8];
    char message[32];

    puts("username:");
    gets(username);
    puts("password:");
    gets(password);
    allowed = strcmp("admin", username) + \
               strcmp("topsecre", password);

    puts("message:");
    gets(message);

    printf("user=%s pass=%s result=%d\n", username, \
           password, allowed);

    if(allowed == 0)
        printf("Access granted. Message sent!\n");
    else
        printf("Access denied\n");

    return 0;
}
```

Classic Overflow – classic/prog_2

- Variable order will determine how it can be exploited
 - Implementation dependent
- message is the prime suspect as it is written after the evaluation is done
- Can also change **an internal decision (flow inside the function)** by writing over the allowed variable

```
int main() {
    char allowed = 0;
    char password[8];
    char username[8];
    char message[32];

    puts("username:");
    gets(username);
    puts("password:");
    gets(password);
    allowed = strcmp("admin", username) + \
               strcmp("topsecre", password);

    puts("message:");
    gets(message);

    printf("user=%s pass=%s result=%d\n", username, \
           password, allowed);

    if(allowed == 0)
        printf("Access granted. Message sent!\n");
    else
        printf("Access denied\n");

    return 0;
}
```

Classic Overflow – classic/prog_2

p &allowed

0x7fffffffedf2f

p &username

0x7fffffffedf1f

p &password

0x7fffffffedf27

p &message

0x7fffffffedef0

Memory grows from top to bottom

message can be used to overwrite everything!!!

allowed

```
0x00007fffffffedef0 | +0x0000: 0x006567617373656d ("message"? ) ← $rax, $rsp, $r8
0x00007fffffffedef8 | +0x0008: 0x00007fffffffedf27 → 0x64726f7773736170
0x00007fffffffedef00 | +0x0010: 0x00007fffffffedf26 → 0x726f777373617000
0x00007fffffffedef08 | +0x0018: 0x0000000000004012cd → <__libc_csu_init+77> add rbx, 0x1
0x00007fffffffedef10 | +0x0020: 0x00007fffffff790fc8 → 0x0000000000000000
0x00007fffffffedef18 | +0x0028: 0x610000000000401280
0x00007fffffffedef20 | +0x0030: 0x700000000006e696d64 ("dmin"? )
0x00007fffffffedef28 | +0x0038: 0x0464726f77737361
0x00007fffffffedef30 | +0x0040: 0x00007fffffffefee030 → 0x0000000000000001
0x00007fffffffedef38 | +0x0048: 0x000000000000401280 → <__libc_csu_init+0> endbr64
0x00007fffffffedef40 | +0x0050: 0x000000000000000000 ← $rbp
0x00007fffffffedef48 | +0x0058: 0x00007fffffff5c70b3 → <__libc_start_main+243> mov edi, eax
```

message

username

password

Exercise: classic/prog 2

- Compile the binary: `gcc -g -O0 -fno-stack-protector -o prog_2 prog_2.c`
- Analyze the execution with different payloads
 - Print register: `p $rsp` or variable address `p &username`
 - Check stack information: `info frame`
- Determine
 - What is the stack base address?
 - Where is the return information?
 - How many bytes can be entered to the message without overflow?
 - How many bytes can be written without damage?
 - What happens when an overflow is achieved?
 - How can the decision be subverted?

CWE-126: Buffer Over-read

- **The software reads from a buffer and reference memory locations after the targeted buffer.**
 - using buffer access mechanisms such as indexes or pointers
- **Impact: Allows access to otherwise private data**
- **Most common with:**
 - Casts between structures with different sizes
 - Copy of data without considering the actual size, assuming a general size
 - Copy of data based on corrupted metadata
 - Erasure of `\0` in null terminated strings

Buffer Over-read – overread1.c

- **Program flow:**
 - Program reads a string without boundary checks
 - Memory is manipulated
 - A message is printed
- **Demonstrates a read beyond bounds with printf**
- **Impact: private data (message) is disclosed to users**

```
int main(int argc, char* argv[]){
    char message[32];
    char buffer[8];

    printf("Password: ");
    gets(buffer);

    sprintf(message, "Secret message");

    if(strcmp(buffer, "password") == 0) {
        printf("%s\n", message);
    }else{
        printf("Password %s is incorrect\n", buffer);
    }
}
```

Buffer Over-read – overread1

➤ Vulnerability:

- In some situations, the password may overflow the buffer, and further memory operations erase the `\0` character
- Further **printf** of a message will include additional memory

```
int main(int argc, char* argv[]){
    char message[32];
    char buffer[8];

    printf("Password: ");
    gets(buffer);


    sprintf(message, "Secret message");

    if(strcmp(buffer, "password") == 0) {
        printf("%s\n", message);
    }else{
        printf("Password %s is incorrect\n", buffer);
    }
}
```

Buffer Over-read – overread1

➤ Exercise: Determine what conditions trigger the vulnerability, and what is the impact.

➤ Write overflow



➤ Memory manipulation erase end of string (`\0`)



➤ Read overflow



Buffer Over-read – server.c

➤ Program Flow

- Receives a message to a buffer
- Prints the buffer
- Returns the buffer through the socket

➤ Vulnerability:

- Send doesn't respect buffer sizes and will use a buffer larger than expected
- **printf** has no notion of string size and will print everything up to `\0`

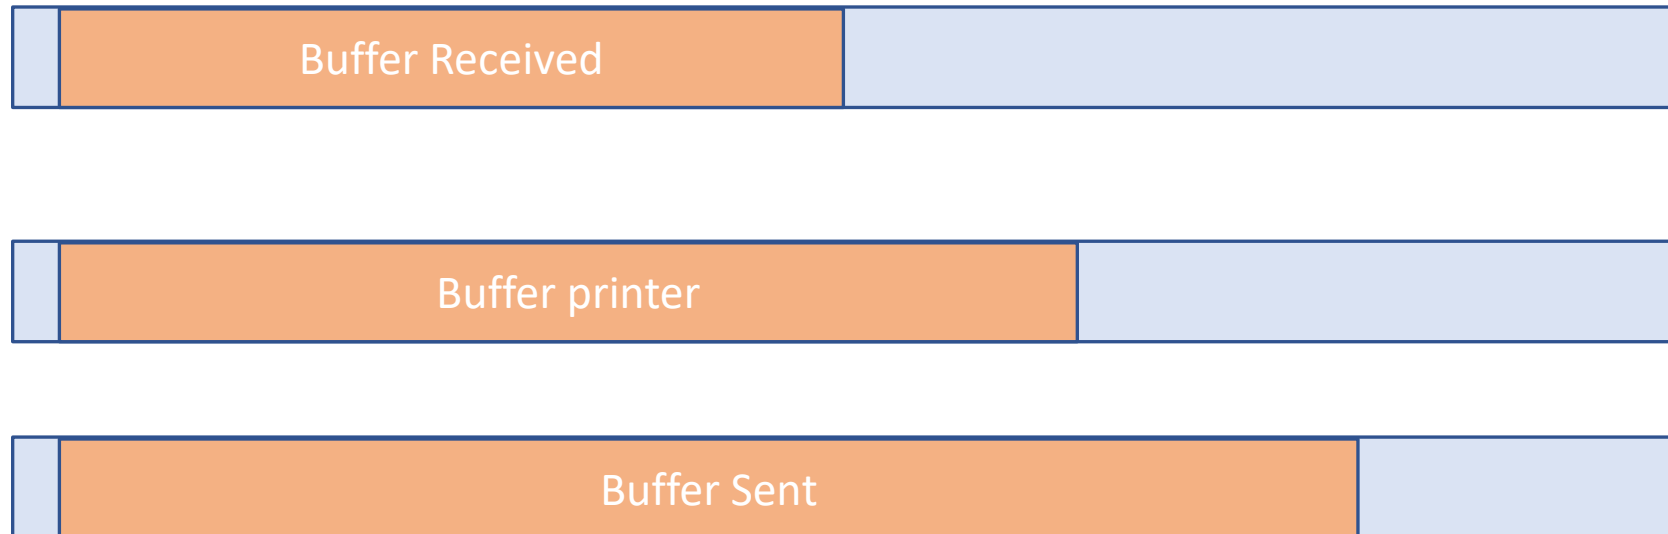
➤ Impact: existing memory contents will be sent to clients

```
while(1){  
    n = recvfrom(sockfd, buffer, 32, NULL, &cliaddr, &len);  
    printf("%s\n", buffer);  
    sendto(sockfd, buffer, MESSAGE_SIZE, NULL, &cliaddr, len);  
}
```

Buffer Over-read – server.c

➤ **Exercise: Determine what conditions trigger the vulnerability, and what is the impact.**

➤ **Variable structure:**



Stack Overflow


Stack Based Vulnerabilities

➤ Stack can be subverted to conduct attacks

- it contains local variables (which store user injected data)
- the program execution flow is kept in the stack

➤ Mostly:

- Denial of Service: program crashes
- Memory disclosure: attacker gains access to previous frames
- Change program flow
- Injection of malicious code



Stack Based Vulnerabilities

➤ Recap...

➤ Local variables will overwrite others

- Can change data stored
- Can lead to local memory disclosure
- Can change local decisions if they depend of stored data



Stack Based Vulnerabilities

➤ Recap...


➤ Local variables will overwrite others

- Can change data stored
- Can lead to local memory disclosure
- Can change local decisions if they depend of stored data

➤ Further writing will overwrite flow information

- If done blindly, program will crash (why?)

➤ It affects frames from previous functions



Stack Smashing

- **What about writing the correct values to the stack?**
 - Some value to RBP
 - An address belonging to the process in RIP
- **Well... when the message ends the flow will be restored**
 - That is... stored RBP and stored RIP are loaded into the registers
 - The stack frame will start at RBP
 - Program jump to the address in RIP
- **If the addresses aren't in a mapped area, program will receive a SIGSEV**



Stack: program_flow.c

➤ Program flow:

- Reads data from file
- Calls **foo** function with size and buffer
- **foo** has an overflowing **memcpy**
- **secret** function is never called

➤ Attack: Overflow the buffer

- writing over stored \$RBP
- writing over stored \$RIP, placing **&secret** there

➤ Consider ASLR to be disabled

```
void secret(){
    printf("Secret message\n");
    exit(0);
}

char foo(int size, char* arg){
    char buffer[8];
    memcpy(buffer, arg, size);
    return buffer[0];
}

int main(int argc, char* argv[]){
    char buffer[64];
    printf("%p\n", &secret);

    FILE *fp = fopen(argv[1], "r");
    int size = fread(buffer, 1, 64, fp);
    fclose(fp);

    foo(size, buffer);
    return 0;
}
```

Stack: program_flow.c

➤ **Main stack** →

➤ **Foo stack** →

- Stored program flow

- **buffer[8]** →

➤ **Secret has no stack!**



```
void secret(){
    printf("Secret message\n");
    exit(0);
}

char foo(int size, char* arg){
    char buffer[8];
    memcpy(buffer, arg, size);
    return buffer[0];
}

int main(int argc, char* argv[]){
    char buffer[64];
    printf("%p\n", &secret);

    FILE *fp = fopen(argv[1], "r");
    int size = fread(buffer, 1, 64, fp);
    fclose(fp);

    foo(size, buffer);
    return 0;
}
```


Stack: program_flow.c

➤ Attack strategy

- Overwrite buffer over RBP/RIP

➤ How to find the addresses?

- If we have the source code:
printf("%p\n", secret);
- If we don't: **gdb** or bruteforce



```
void secret(){
    printf("Secret message\n");
    exit(0);
}

char foo(int size, char* arg){
    char buffer[8];
    memcpy(buffer, arg, size);
    return buffer[0];
}

int main(int argc, char* argv[]){
    char buffer[64];
    printf("%p\n", &secret);

    FILE *fp = fopen(argv[1], "r");
    int size = fread(buffer, 1, 64, fp);
    fclose(fp);

    foo(size, buffer);
    return 0;
}
```

Stack: program_flow.c

```
$ ./program_flow payload
```

```
0x8001209
```

Value to inject
program vs **gdb**
may yield different
values!

```
$ gdb program_flow payload
```

```
gdb$ br main
```

```
gdb$ run
```

```
gdb$ print &secret
```

```
gdb$ 5 = (void (*)( )) 0x8001209  
<secret>
```

```
void secret(){  
    printf("Secret message\n");  
    exit(0);  
}  
char foo(int size, char* arg){  
    char buffer[8];  
    memcpy(buffer, arg, size);  
    return buffer[0];  
}  
int main(int argc, char* argv[]){  
    char buffer[64];  
    printf("%p\n", &secret);  
  
    FILE *fp = fopen(argv[1], "r");  
    int size = fread(buffer, 1, 64, fp);  
    fclose(fp);  
  
    foo(size, buffer);  
    return 0;  
}
```

Stack Smashing – program_flow.c

➤ Typical flow



```
void secret(){
    printf("Secret message\n");
    exit(0);
}
char foo(int size, char* arg){
    char buffer[8];
    memcpy(buffer, arg, size);
    return buffer[0];
}
int main(int argc, char* argv[]){
    char buffer[64];
    printf("%p\n", &secret);

    FILE *fp = fopen(argv[1], "r");
    int size = fread(buffer, 1, 64, fp);
    fclose(fp);

    foo(size, buffer);
    return 0;
}
```


Stack: program_flow.c

➤ Flow subverted to secret()



```
void secret(){
    printf("Secret message\n");
    exit(0);
}

char foo(int size, char* arg){
    char buffer[8];
    memcpy(buffer, arg, size);
    return buffer[0];
}

int main(int argc, char* argv[]){
    char buffer[64];
    printf("%p\n", &secret);

    FILE *fp = fopen(argv[1], "r");
    int size = fread(buffer, 1, 64, fp);
    fclose(fp);

    foo(size, buffer);
    return 0;
}
```

Stack: program_flow.c

```
$ program_flow payload
```

```
0x8001209
```

```
Secret message
```

With the correct
payload, secret() is
called

Q: What payload?

```
void secret(){
    printf("Secret message\n");
    exit(0);
}

char foo(int size, char* arg){
    char buffer[8];
    memcpy(buffer, arg, size);
    return buffer[0];
}


int main(int argc, char* argv[]){
    char buffer[64];
    printf("%p\n", &secret);

    FILE *fp = fopen(argv[1], "r");
    int size = fread(buffer, 1, 64, fp);
    e(fp);

    size, buffer);
    n 0;
```

Stack: return_to_libc.c

- **Instead of returning to a program function it is possible to jump to other locations**
 - In theory, any segment allocated to the program
 - In practice, permission mechanisms limit the available segments
- **Segments for libraries have several generic libraries**
 - In particular: `system()`
 - Is mostly executable
- **Stack can be executable**
 - but it isn't on recent systems



Stack: return_to_libc.c

➤ Typical Flow



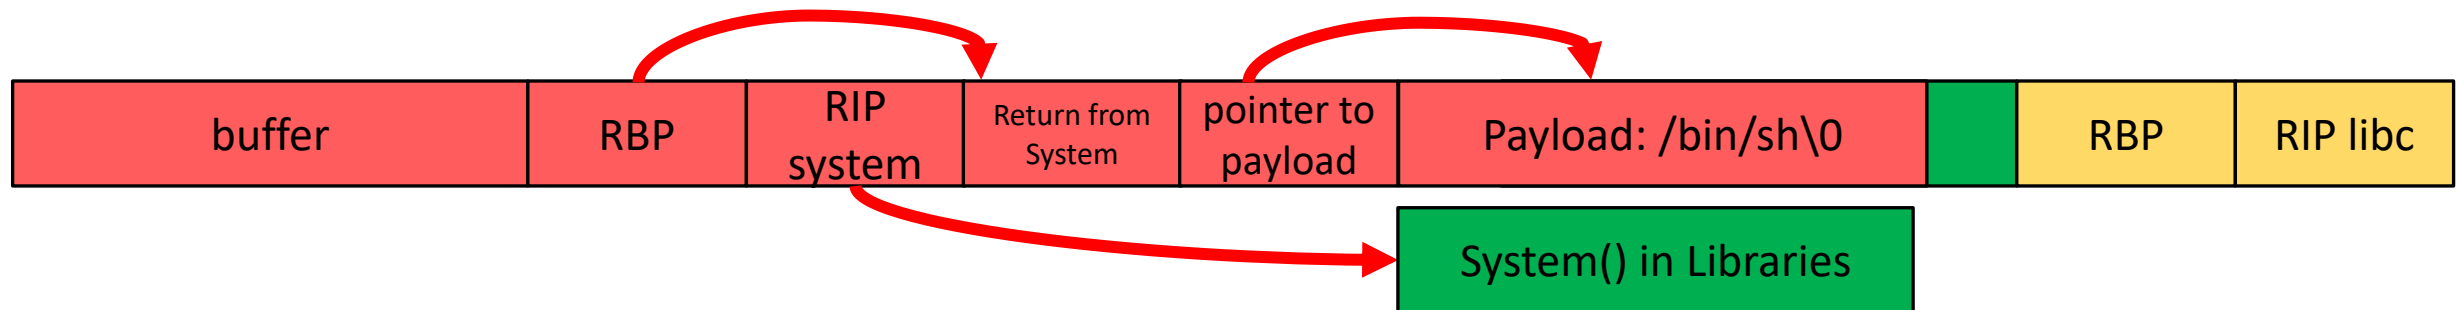
➤ Return to libc

- Build “fake” Stack frame and call `system()` with one argument
 - Argument is the command to execute (e.g. a reverse shell)
- Must take in consideration calling convention
 - Which is architecture dependent

Stack: return_to_libc.c (32bits)

➤ Arguments are passed in the stack

- Approach: store values to the stack so that system is called with a payload
 - Then call system



Countermeasures: Data Executable Prevention

➤ **Non Executable Stack (NX) (Data Executable Prevention)**

- Most binaries do not allow running code from Stack
- Stack segments are marked as Non Executable (NX bit)
 - code cannot jump to it
 - Return to lib-c attack not possible

➤ **Introduced in recent OS, but can be disabled**

- Not ubiquitous on embedded devices
- Binaries must opt-in!

Countermeasures: Canaries

➤ Uses references values after local variables to detect overflow

- Value is placed when the function starts
- Value is compared before function exits
- Program is interrupted if values do not match

➤ Stack canaries:



Countermeasures: Canaries

Without Canaries

```
push    rbp
mov     rbp, rsp
sub     rsp, 16
lea     rax, -10[rbp]
mov     rsi, rax
lea     rdi, .LC0[rip]
mov     eax, 0
call    __isoc99_scanf@PLT
lea     rax, -10[rbp]
mov     rdi, rax
call    puts@PLT
nop
leave
ret
```

With Canaries

```
push    rbp
mov     rbp, rsp
sub     rsp, 32
mov     rax, QWORD PTR fs:40
mov     QWORD PTR -8[rbp], rax
xor     eax, eax
lea     rax, -18[rbp]
mov     rsi, rax
lea     rdi, .LC0[rip]
mov     eax, 0
call    __isoc99_scanf@PLT
lea     rax, -18[rbp]
mov     rdi, rax
call    puts@PLT
nop
mov     rax, QWORD PTR -8[rbp]
xor     rax, QWORD PTR fs:40
je      .L2
call    __stack_chk_fail@PLT
```

L2:

```
leave
ret
```

Gets value from fs:40
Stores value at rbp-8 (inside
stack frame)

Fetches value
Xor with reference at fs:40
Exit or crash

Countermeasures: Canaries

- **-fno-stack-protector:** disables stack protection. (What we have been using)
- **-fstack-protector:** enables stack protection for vulnerable functions that contain:
 - A character array larger than 8 bytes.
 - An 8-bit integer array larger than 8 bytes.
 - A call to `alloca()` with either a variable size or a constant size bigger than 8 bytes.
- **-fstack-protector-strong:** enables stack protection for vulnerable functions that contain:
 - An array of any size and type.
 - A call to `alloca()`.
 - A local variable that has its address taken.
- **-fstack-protector-all:** adds stack protection to all functions regardless of their vulnerability.

Stack: return_to_libc.c (x86_64)

➤ **x64: first arguments are passed in register: RDI, RSI, RDX, RCX**

- Approach: load RDI with address of string, jump to system address
- Problems: cannot jump to stack (due to NX)

➤ **Improved:**

- Search any code that loads RDI from stack
 - we can control what is in the stack but we cannot execute code from it
- jump to code that loads RDI from stack
- Jump to system

ROP

- **Return Oriented Programming: Execute code already present in the program.**
 - Each snippet is composed by some instructions + **RET**
 - **RET** pops RIP from the stack
- **Program flow is controlled by values in the stack**
 - Attacker puts values in stack pointing to gadgets
 - When a gadget ends, the code jumps to the next gadget
- **Any program can be constructed as long as there are gadgets available**
 - When Good Instructions Go Bad: Generalizing Return-Oriented Programming to RISC [1] - Buchanan, E.; Roemer, R.; Shacham, H.; Savage, S.
 - Return-Oriented Programming: Exploits Without Code Injection [2] - Shacham, Hovav; Buchanan, Erik; Roemer, Ryan; Savage, Stefan.

ROP

- **ROP Attacks: Chain gadgets to execute malicious code.**
- **A gadget is a suite of instructions which end by the branch instruction `ret` (Intel) or the equivalent on ARM.**

Intel examples:

```
pop eax ; ret  
xor ebx, ebx ; ret
```

ARM examples:

```
pop {r4, pc}  
str r1, [r0] ; bx lr
```

- **Objective: Use gadgets instead of classical shellcode**

ROP

- **Because x86 instructions aren't aligned, a gadget can contain another gadget.**

```
f7c707000000f9545c3 → test edi, 0x7 ; setnz byte ptr [rbp-0x3d] ;  
c70700000000f9545c3 → mov dword ptr [rdi], 0xf00000 ; xchg ebp, eax ; ret
```

- **Doesn't work on RISC architectures like ARM, MIPS, SPARC...**

```

0x00000000000040124c: mov rsi, rcx; mov rdi, rax; call 0x10e0; movzx eax, byte ptr [rbp - 8]; leave; ret;
0x000000000000401306: mov rsi, rdx; mov rdi, rax; call 0x1214; mov eax, 0; leave; ret;
0x000000000000401257: movzx eax, byte ptr [rbp - 8]; leave; ret;
0x000000000000401388: nop dword ptr [rax + rax]; endbr64; ret;
0x000000000000401387: nop dword ptr cs:[rax + rax]; endbr64; ret;
0x000000000000401386: nop word ptr cs:[rax + rax]; endbr64; ret;
0x000000000000401007: or byte ptr [rax - 0x75], cl; add eax, 0x2fe9; test rax, rax; je 0x1016; call rax;
0x000000000000401166: or dword ptr [rdi + 0x404060], edi; jmp rax;
0x00000000000040137c: pop r12; pop r13; pop r14; pop r15; ret;
0x00000000000040137e: pop r13; pop r14; pop r15; ret;
0x000000000000401380: pop r14; pop r15; ret;
0x000000000000401382: pop r15; ret;
0x00000000000040137b: pop rbp; pop r12; pop r13; pop r14; pop r15; ret;
0x00000000000040137f: pop rbp; pop r14; pop r15; ret;
0x0000000000004011dd: pop rbp; ret;
0x000000000000401383: pop rdi; ret;
0x000000000000401381: pop rsi; pop r15; ret;
0x00000000000040137d: pop rsp; pop r13; pop r14; pop r15; ret;
0x0000000000004011cd: push rbp; mov rbp, rsp; call 0x1150; mov byte ptr [rip + 0x2e83], 1; pop rbp; ret;
0x0000000000004012ea: ret 0xffffd;
0x000000000000401011: sal byte ptr [rdx + rax - 1], 0xd0; add rsp, 8; ret;
0x0000000000004011d8: sub dword ptr [rsi], 0; add byte ptr [rcx], al; pop rbp; ret;
0x00000000000040139d: sub esp, 8; add rsp, 8; ret;
0x000000000000401005: sub esp, 8; mov rax, qword ptr [rip + 0x2fe9]; test rax, rax; je 0x1016; call rax;
0x00000000000040139c: sub rsp, 8; add rsp, 8; ret;
0x000000000000401004: sub rsp, 8; mov rax, qword ptr [rip + 0x2fe9]; test rax, rax; je 0x1016; call rax;
0x00000000000040138a: test byte ptr [rax], al; add byte ptr [rax], al; add byte ptr [rax], al; endbr64; ret;
0x000000000000401010: test eax, eax; je 0x1016; call rax;
0x000000000000401010: test eax, eax; je 0x1016; call rax; add rsp, 8; ret;
0x000000000000401163: test eax, eax; je 0x1170; mov edi, 0x404060; jmp rax;
0x0000000000004011a5: test eax, eax; je 0x11b0; mov edi, 0x404060; jmp rax;
0x00000000000040100f: test rax, rax; je 0x1016; call rax;
0x00000000000040100f: test rax, rax; je 0x1016; call rax; add rsp, 8; ret;
0x000000000000401162: test rax, rax; je 0x1170; mov edi, 0x404060; jmp rax;
0x0000000000004011a4: test rax, rax; je 0x11b0; mov edi, 0x404060; jmp rax;

```

ROP

➤ Using ROP, stack is subverted to create a jump sequence. It contains:


- Values to be loaded
- Addresses to other gadgets
- May also contain arguments to functions called

➤ Gadgets are present in program code and loaded libraries

- Each function available provides one gadget
- Plus misaligned access

➤ Why?

- It can bypass several security mechanisms




```

0x000000000000401011: sal byte ptr [rdx + rax - 1], 0xd0; add rsp, 8; ret;
0x0000000000004011d8: sub dword ptr [rsi], 0; add byte ptr [rcx], al; pop rbp; ret;
0x00000000000040139d: sub esp, 8; add rsp, 8; ret;
0x000000000000401005: sub esp, 8; mov rax, qword ptr [rip + 0x2fe9]; test rax, rax; je 0x1016; call rax;
0x00000000000040139c: sub rsp, 8; add rsp, 8; ret;
0x000000000000401004: sub rsp, 8; mov rax, qword ptr [rip + 0x2fe9]; test rax, rax; je 0x1016; call rax;
0x00000000000040138a: test byte ptr [rax], al; add byte ptr [rax], al; add byte ptr [rax], al; endbr64; ret;
0x000000000000401010: test eax, eax; je 0x1016; call rax;
0x000000000000401010: test eax, eax; je 0x1016; call rax; add rsp, 8; ret;
0x000000000000401163: test eax, eax; je 0x1170; mov edi, 0x404060; jmp rax;
0x0000000000004011a5: test eax, eax; je 0x11b0; mov edi, 0x404060; jmp rax;
0x00000000000040100f: test rax, rax; je 0x1016; call rax;
0x00000000000040100f: test rax, rax; je 0x1016; call rax; add rsp, 8; ret;
0x000000000000401162: test rax, rax; je 0x1170; mov edi, 0x404060; jmp rax;
0x0000000000004011a4: test rax, rax; je 0x11b0; mov edi, 0x404060; jmp rax;
0x00000000000040125a: cld; leave; ret;
0x00000000000040139b: cli; sub rsp, 8; add rsp, 8; ret;
0x000000000000401003: cli; sub rsp, 8; mov rax, qword ptr [rip + 0x2fe9]; test rax, rax; je 0x1016; call rax;
0x000000000000401143: cli; ret;
0x000000000000401398: endbr64; sub rsp, 8; add rsp, 8; ret;
0x000000000000401000: endbr64; sub rsp, 8; mov rax, qword ptr [rip + 0x2fe9]; test rax, rax; je 0x1016; call rax;
0x000000000000401140: endbr64; ret;
0x00000000000040113e: hlt; nop; endbr64; ret;
0x00000000000040125b: leave; ret;
0x00000000000040113f: nop; endbr64; ret;
0x00000000000040116f: nop; ret;
0x00000000000040101a: ret;

```

111 gadgets found

```

gef> rop --search 'pop rdi'
[INFO] Load gadgets from cache
[LOAD] loading... 100%
[LOAD] removing double gadgets... 100%
[INFO] Searching for gadgets: pop rdi


```

```

0x000000000000401383: pop rdi; ret;

```


Stack: return_to_libc.c (x86_64)




➤ Payload strategy:

- All addresses are 8 bytes
- Buffer: padding with 16 bytes (buffer + RBP)
- Gadget address: ?? -> **rop --search "pop rdi; ret"**
 - pop RDI: load command address into RDI
 - ret: load system address into RIP
- Command address: ?? -> **grep /bin/sh**
 - Approaches: **Find a string already in RAM (better)**; add the payload after the system address (if required)
- System address: ?? -> **print system**

```

0x00000000000040100f: test rax, rax; je 0x1016; call rax;
0x00000000000040100f: test rax, rax; je 0x1016; call rax; add rsp, 8; ret;
0x000000000000401162: test rax, rax; je 0x1170; mov edi, 0x404060; jmp rax;
0x0000000000004011a4: test rax, rax; je 0x11b0; mov edi, 0x404060; jmp rax;
0x00000000000040125a: clc; leave; ret;
0x00000000000040139b: cli; sub rsp, 8; add rsp, 8; ret;
0x000000000000401003: cli; sub rsp, 8; mov rax, qword ptr [rip + 0x2fe9]; test rax, rax; je 0x1016; call rax;
0x000000000000401143: cli; ret;
0x000000000000401398: endbr64; sub rsp, 8; add rsp, 8; ret;
0x000000000000401000: endbr64; sub rsp, 8; mov rax, qword ptr [rip + 0x2fe9]; test rax, rax; je 0x1016; call rax;
0x000000000000401140: endbr64; ret;
0x00000000000040113e: hlt; nop; endbr64; ret;
0x00000000000040125b: leave; ret;
0x00000000000040113f: nop; endbr64; ret;
0x00000000000040116f: nop; ret;
0x00000000000040101a: ret;

```



111 gadgets found

gef> print system

\$14 = {<text variable, no debug info>} 0x7ffffff5f5410 <system>

gef> grep "/bin/sh"

[+] Searching '/bin/sh' in memory

[+] In '[heap]'(0x405000-0x426000), permission=rw-

0x4058b8 - 0x4058bf → "/bin/sh"

[+] In '/usr/lib/x86_64-linux-gnu/libc-2.31.so'(0x7ffffff73d000-0x7ffffff787000), permission=r--

0x7ffffff7575aa - 0x7ffffff7575b1 → "/bin/sh"

[+] In '0x7ffffff7df000-0x7ffffff7e2000), permission=rw-

0x7ffffff7e1db0 - 0x7ffffff7e1db7 → "/bin/sh"


[+] In '[stack]'(0x7ffffff7ef000-0x7ffffff7fef000), permission=rw-

0x7ffffff7fedf30 - 0x7ffffff7fedf37 → "/bin/sh"

0x7ffffff7fedf58 - 0x7ffffff7fedf5f → "/bin/sh[...]"

gef> |

Stack: return_to_libc.c (x86_64)



➤ Payload strategy:

- All addresses are 8 bytes
- Buffer: padding with 16 bytes (buffer + RBP)
- Gadget address: 0x00401383
 - pop RDI: load command address into RDI
 - ret: load system address into RIP
- Command address: 0x7fffffff7575aa
 - Approaches: **Find a string already in RAM (better)**; add the payload after the system address (if required)
- System address: 0x7fffffff5f5410

Stack: return_to_libc.c (x86_64)

buffer	RBP	Gadget1 address	Gadget2 address	Command address	System address	Command to be executed (optional)	RBP	RIP libc
--------	-----	--------------------	--------------------	--------------------	-------------------	--------------------------------------	-----	----------

➤ In some systems, stack must be aligned to 16 bytes and our ROP chain isn't...

- Result is a crash in instruction `movaps`

➤ **Solution: add another gadget with only a `ret` (will pop a value)**

- Gadget 1: `0x00401384 ; ret`
- Gadget 2: `0x00401383 ; pop rdi;ret`

Stack: return_to_libc.c (x86_64)

➤ Exercise: build a ROP chain and get a shell in the program

- It may be useful to disable ASLR for now
 - In gef: `aslr off`
 - System wide (as root): `echo 0 > /proc/sys/kernel/randomize_va_space`
- Document the payload

➤ Exercise: build a ROP chain to start a remote shell

- Document the payload and the differences from the previous

ROP Variants

➤ JOP: Jump Oriented Programming

- <https://www.comp.nus.edu.sg/~liangzk/papers/asiaccs11.pdf>

➤ SOP: **Jump** Oriented Programming

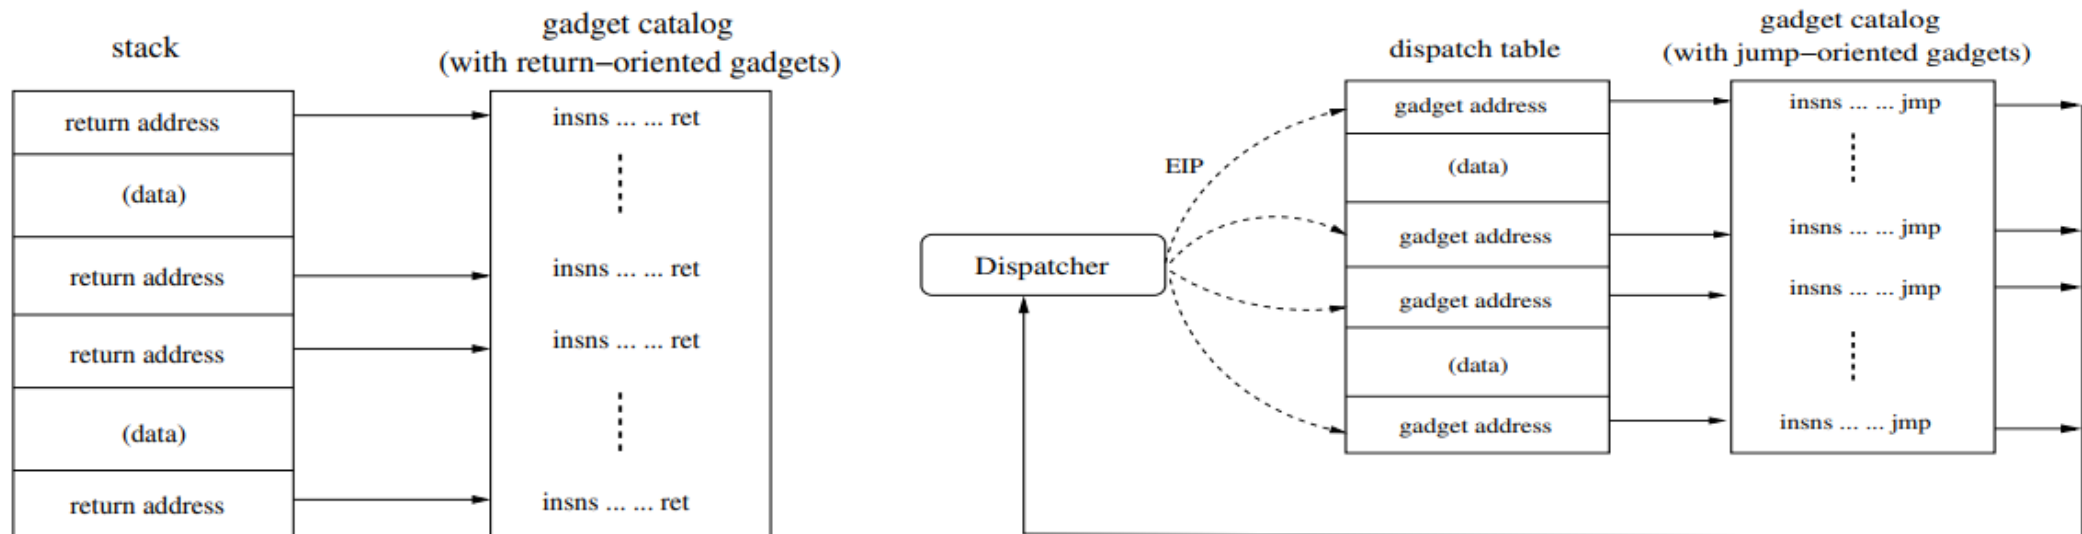
- https://www.lst.inf.ethz.ch/research/publications/PPREW_2013/PPREW_2013.pdf

➤ BROP: Blind Return Oriented Programming

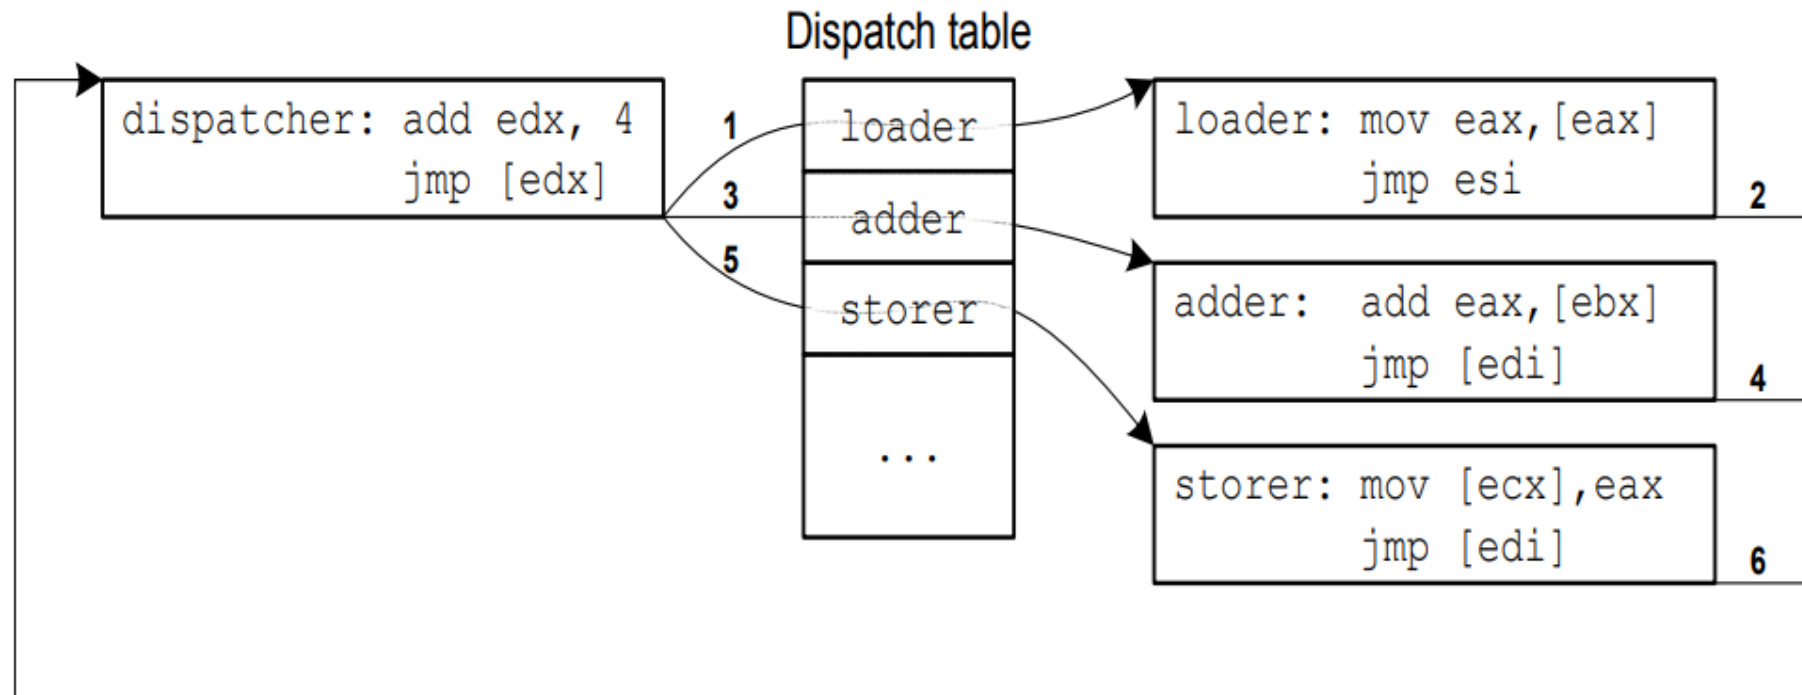
- <http://www.scs.stanford.edu/brop/bittau-brop.pdf>

Jump Oriented Programming

- Explores small gadgets that end with an indirect JMP with a dispatcher
 - Indirect jmp: jmp [register]
 - Is assumed to be more complex to detect and avoid as interaction is restricted to code and registers
 - Although number of JMP gadgets is smaller, unaligned execution **create jumps** not previously present in the code
 - The program counter is **any register**



Jump Oriented Programming



String Oriented Programming

➤ Makes use of a String format bug

- Present in the printf family of functions (printf, vprintf, fprintf)
- Correct: printf("%s", str);
- Vulnerable: printf(str);

➤ Format string attacks read/write arbitrary values to arbitrary memory locations

- Explore %p, %n, %s,
- Can be used to trigger ROP, JOP attacks by writing values memory
- Instead of writing sequential chunks, SOP can issue arbitrary writes.

➤ Two approaches

- Direct control flow redirect: Erase return value on stack, jumping to gadget on function end
- Indirect control flow redirect: Erase a Global Offset Table entry
 - GOT keeps addresses to external symbols as resolved by the linker

Blind Return Oriented Programming


- **Makes it possible to write exploits without possessing the target's binary.**
 - It requires a stack overflow and a service that restarts after a crash.
 - Based on whether a service crashes
 - Is able to construct a full remote exploit that leads to a shell.
- **The attack remotely leaks gadgets to perform the write system call, after which the binary is transferred from memory to the attacker's socket.**
 - Following that, a standard ROP attack can be carried out.
 - Apart from attacking proprietary services, BROP is very useful in targeting open-source software for which the particular binary used is not public

Blind Return Oriented Programming

- **Makes it possible to write exploits without possessing the target's binary.**
 - It requires a stack overflow and a service that restarts after a crash.
 - Based on whether a service crashes
 - Is able to construct a full remote exploit that leads to a shell.
- **The attack remotely leaks gadgets to perform the write system call, after which the binary is transferred from memory to the attacker's socket.**
 - Following that, a standard ROP attack can be carried out.
 - Apart from attacking proprietary services, BRPOP is very useful in targeting open-source software for which the particular binary used is not public

Blind Return Oriented Programming

- Looks for specific ROP Gadgets until a specific combination is found



Blind Return Oriented Programming

➤ The BROP attack has the following phases:

1. Stack reading: read the stack to leak canaries and a return address to defeat ASLR.

Method: overflows varying the last byte. Byte found if app doesn't crash
512-640 requests required

2. Blind ROP: find enough gadgets to invoke write and control its arguments.

Method: find a Gadget1 that stops the service. Then brute force other gadgets together with this.

Implement a clever method to identify different gadgets

3. Build the exploit: dump enough of the binary to find enough gadgets to build a shellcode, and launch the final exploit.

Obtain access to the write call so that the binary can be dumped

Heap Overflow

Heap Overflow

- **Heap is used to store dynamically allocated variables**
 - Allocation: malloc, calloc and new (C++), release: free or delete (C++)
- **Call reserves a chunk and returns a pointer to the buffer**
 - buffer: $(8 + (n / 8) * 8)$ bytes
 - If chunk is free data will have
 - Forward Pointer (4 bytes), pointer to next free chunk
 - Backwards Pointer (4 bytes), pointer to previous free chunk
 - Headers used for housekeeping
 - Previous Chunk Size (previous chunk is free), 4 bytes
 - Chunk Size + flags, 4 bytes
 - Flags
 - 0x01 PREV_INUSE – set when previous chunk is in use
 - 0x02 IS_MMAPPED – set if chunk was obtained with mmap()
 - 0x04 NON_MAIN_ARENA – set if chunk belongs to a thread arena



Heap Overflow: overflow.c

➤ Overflow/underflow will write/read over control structures and then data

- Control structures are implementation specific
- As well as reuse and actual buffer location

```
int main(int argc, char **argv) {
    char *buf1 = (char *) malloc(BUFSIZE);
    char *buf2 = (char *) malloc(BUFSIZE);
    memset(buf1, 0, BUFSIZE); //Clear data
    memset(buf2, 0, BUFSIZE);

    printf("Buf2: %s\n", buf2); //Should print "Buf2: "
    strcpy(buf1, argv[1]);
    printf("Buf2: %s\n", buf2); //Should print "Buf2: "
}
```

prev size

size

buffer

prev size

size

buffer

prev size

size

buffer

Heap Overflow: dangling.c

➤ Dangling references can give access to memory

- Both for read and write purposes

```
char *buf1 = (char *) malloc(BUFSIZE*100); //Allocate buffer
memset(buf1, 'U', BUFSIZE); //Fill it with 0x55
free(buf1); //Free the memory
```

```
char *buf2 = (char *) malloc(BUFSIZE); //Allocate new buffer
memset(buf2, 'A', BUFSIZE); //Fill it with 0x41
```

```
printf("%s\n", buf1); //buf1 was freed
```

➤ Access to buf1 should be denied: it isn't

➤ Access to buf1 should not give access to other ranges: it gives to buf2

prev size

size

buffer

prev size

size

buffer

prev size

size

buffer

Heap Overflow: fastbin.c

➤ Glibc has lists of recently freed blocks

- Each list (bin) stores chunks with a specific size
- Blocks are reused in future allocations if size is compatible
 - Great for performance as the memory is already reserved
 - Horrible for security as dangling pointers will give a view to memory areas

➤ Bins are also used to detect double free

- We cannot free a chunk that rests at the top of the bin
- Which is great for security as a double free could corrupt the linked list

Heap Overflow: fastbin.c

- **Fast Bin attack explores Bins to get a pointer to an already allocated area**
 - Result is program will have **two pointers to the same memory**
 - Especially useful if memory stores dynamic objects with function, as function pointers can be overwritten
 - The first pointer is legitimate
 - The second is a shadow pointer

- **Attack strategy**
 - Allocate at least three buffers (a, b, c) with the same size
 - To use same bin
 - free(a), then free(b), then free(a) again
 - Double freeing a will ensure that the fast bin will have duplicated entries (a)
 - Bin will have three pointers ready to use: a b a
 - Allocate three buffers again with the same size.
 - Result is a legitimate pointer, another legitimate pointer, and a shadow pointer

Heap Overflow: fastbin.c

➤ Impact: attacker can gain access to memory region

- If victim has chunk a with data and leaks
- Attacker can fill free list and allocate again

```
// Allocating 3 buffers
int *a = calloc(1, 8);
int *b = calloc(1, 8);
int *c = calloc(1, 8);

free(a);
free(b);
free(a); //AGAIN!

//Free list now has: a b a

int *d = calloc(1, 8);
int *e = calloc(1, 8);
int *f = calloc(1, 8);

// d will be equal to f
```

Heap Overflow: overflow.c

- **Exercise: Observe and document the behavior in both programs**
 - dangling.c and overflow.c
 - Use GDB to analyse the addresses
 - What is the impact of writing to a freed pointer?

Countermeasures: ASLR

➤ Address Space Layout Randomization (ASLR)

- Address are dynamic across process execution
 - Different architectures and configurations apply randomization to different segments
 - Only Stack is randomized, all segments are randomized
- Not trivial to predict the address to issue a jump or change memory

➤ `echo $n > /proc/sys/kernel/randomize_va_space`

- 0 = No randomization
- 1 = Conservative Randomization: Stack, Heap, Shared Libs
- 2 = Full Randomization: 1 + memory managed via `brk()`

Effects of ASLR (WSL1 on Windows 10)

➤ randomize_va_space =2

```
main: 0x7f80def82189, argc: 0x7ffffbfce569c, local: 0x7ffffbfce56ac, heap: 0x7ffffb8c4b2a0, libc: 0x7f80ded85410  
main: 0x7fb811d47189, argc: 0x7ffffdbd2928c, local: 0x7ffffdbd2929c, heap: 0x7ffffd47952a0, libc: 0x7fb811b55410  
main: 0x7f95178f0189, argc: 0x7ffffee962b7c, local: 0x7ffffee962b8c, heap: 0x7ffffe67082a0, libc: 0x7f95176f5410
```

➤ randomize_va_space =1

```
main: 0x7f1672f77189, argc: 0x7ffffe5835f0c, local: 0x7ffffe5835f1c, heap: 0x7f1672f7b2a0, libc: 0x7f1672d85410  
main: 0x7f6f0aed0189, argc: 0x7ffffd8eb4e9c, local: 0x7ffffd8eb4eac, heap: 0x7f6f0aed42a0, libc: 0x7f6f0acd5410  
main: 0x7f8106545189, argc: 0x7fffff8601bdc, local: 0x7fffff8601bec, heap: 0x7f81065492a0, libc: 0x7f8106355410
```

➤ randomize_va_space=0

```
main: 0x8001189, argc: 0x7fffffffec, local: 0x7fffffffecfc, heap: 0x80052a0, libc: 0x7fffffff5f5410  
main: 0x8001189, argc: 0x7fffffffec, local: 0x7fffffffecfc, heap: 0x80052a0, libc: 0x7fffffff5f5410  
main: 0x8001189, argc: 0x7fffffffec, local: 0x7fffffffecfc, heap: 0x80052a0, libc: 0x7fffffff5f5410
```

Coutermeasures: PIE

➤ Position Independent Executables

- Executables compiled such that their base address does not matter, 'position independent code'

➤ PIE fully enables ASLR as code can be placed dynamically

- Must be enabled at compile time!!
 - `gcc -pie -fPIE`

➤ Breaking ASLR and PIE: Find a reference to some known function

- Because while addresses change, the change keeps relative distance
- e.g.: if we know `printf` is at `0xbf00332`, we will know where is `system`.

ASLR and relative offsets

main: 0x7f80def82189, argc: 0x7fffbfce569c
main: 0x7fb811d47189, argc: 0x7fffd47952a0
main: 0x7f95178f0189, argc: 0x7fffee962b7c

local: 0x7fffbfce56ac, heap: 0x7fffb8c4b2a0
local: 0x7fffd47952a0, heap: 0x7fffd47952a0
local: 0x7fffee962b8c, heap: 0x7fffe67082a0

libc: 0x7f80ded85410
libc: 0x7fb811b55410
libc: 0x7f95176f5410